

Interpretation and Compilation of Languages

Master Programme in Computer Science

Mário Pereira `mjp.pereira@fct.unl.pt`

Nova School of Science and Technology, Portugal

May 5, 2025

Lecture 9

based on lectures by Jean-Christophe Filliâtre and Léon Gondelman
previous editions by João Costa Seco, Luís Caires, and Bernardo Toninho

Today: Compilation of Object-oriented Languages

1. Object layout
2. Dynamic dispatching

Compiling Object Oriented Languages

Let us explain

- how an object is represented
- how a method call is implemented

Let us use `JAVA` as an example (for the moment).

```
class Vehicle {  
    static int start = 10;  
    int position;  
    Vehicle() { position = start; }  
    void move(int d) { position += d; } }
```

```
class Car extends Vehicle {  
    int passengers;  
    void await(Vehicle v) {  
        if (v.position < position)  
            v.move(position - v.position);  
        else  
            move(10); } }
```

```
class Truck extends Vehicle {  
    int load;  
    void move(int d) {  
        if (d <= 55) position += d; else position += 55; } }
```

An object is a heap-allocated block, containing

- its class (and a few other items of information)
- the values of its fields

The value of an object is a pointer to the block.

Key idea: simple inheritance allows us to store the value of some field x at some fixed position in the block: own fields are placed after inherited fields.

Vehicle
position

Car
position
passengers

Truck
position
load

Note the absence of field `start`, which is static and thus allocated elsewhere (for instance in the data segment).

For each field, the compiler knows its position, that is the offset to add to the object pointer.

If for instance field `position` is at offset `+16`, then expression `e.position` is compiled to

```
...                # compile e in %rcx
movl 16(%rcx), %rax # field position at +16
```

this is sound, even if the compiler **only knows the static type** of `e`, which may differ from the dynamic type (the class of the object).

It could even be a sub-class of `Vehicule` that is not yet defined!

Overriding is the ability to redefine a method in a subclass (so that objects in that subclass behave differently).

Example: in class Truck

```
class Truck extends Vehicle {  
    ...  
    void move(int d) { ... }  
}
```

the method `move`, inherited from class `Vehicle`, is **overridden**.

The essence of OO languages lies in **dynamic method call** $e.m(e_1, \dots, e_n)$ (aka dynamic dispatch / message passing).

To do this, we build **class descriptors** containing addresses to method codes (aka **dispatch table**, **vtable**, etc.)

As for class fields, simple inheritance allows us to store the address of (the code of) method m at a fixed offset in this descriptor.

Class descriptors can be allocated in the data segment; each object points to its class descriptor.

```
class Vehicule { void move(int d) {...} }
class Car extends Vehicule { void await(Vehicule v) {...}}
class Truck extends Vehicule { void move(int d) {...} }
```

descr. Vehicule

Vehicule_move

descr. Car

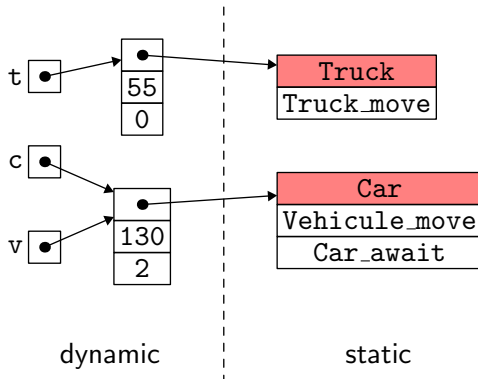
Vehicule_move

Car_await

descr. Truck

Truck_move

```
Truck t = new Truck();
Car c = new Car();
c.passengers = 2;
c.move(60);
Vehicle v = c;
v.move(70);
c.await(t);
```




In practice, the class descriptor for C also points to the class that C inherits from, called the **super class** of C .

This can be a pointer to the descriptor of the super class (for instance stored in the very first slot of the descriptor).


This allows subtyping tests at runtime (*downcast* or *instanceof*).

Class Descriptors, Schematically



```

.data
descr_Vehicle:
    .quad    0
    .quad    Vehicle_move
descr_Car:
    .quad    descr_Vehicle
    .quad    Vehicle_move
    .quad    Car_await
descr_Truck:
    .quad    descr_Vehicle
    .quad    Truck_move
    
```



and the static field **start** of Vehicle is put in the data segment as well

```

static_start:
    .quad    10
    
```

The constructor's code is a function that assumes

- a) the object has been allocated and its address is in %rdi
- b) the first field (class descriptor) is defined and
- c) the arguments are in %rsi, %a2, %rdx, etc., and on the stack.

```
class Vehicle {  
    Vehicle() { position = start; }  
}
```

```
new_Vehicle:  
    movq    $static_start, %rsi  
    movq    %rsi, 8(%rdi)  
    ret
```

For the methods, we adopt the same convention: the object is in %rdi and the arguments of the method in rsi, rdx, etc., and stack if needed.

```
class Vehicle {  
    void move(int d) { position += d; }  
}
```

```
Vehicle_move:  
    addq    %rsi, 8(%rdi)  
    ret
```

(similarly for the move method of Truck)

To compile a call such as `e.move(10)`

1. we compile `e`; its value is a pointer to an object
2. this object contains a pointer to its class descriptor
3. inside, the code for method `move` is located at some offset known from the compiler (for instance `+8`)

```
...           # compile e into %rdi
movq $10, %rsi # parameter
movq (%rdi), %rcx # get the class descriptor
call *8(%rcx)   # call method move
```

As for field access, the compiler has no need to know the actual class of the object (the dynamic type).

If we write

```
Truck v = new Truck();  
((Vehicule)v).move();
```

this is the method `move` from class `Truck` that is called
since the call is always compiled the same way.

The cast only has an influence on the static type
(existence of the method + overloading resolution; see lecture 6).

A Few Words in C++

Let us reuse the vehicles example

```
class Vehicle {  
    static const int start = 10;  
public:  
    int position;  
    Vehicle() { position = start; }  
    virtual void move(int d) { position += d; }  
};
```

virtual means that method `move` can be overridden.

```
class Car : public Vehicle {  
public:  
    int passengers;  
    Car() {}  
    void await(Vehicle &v) { // call by reference  
        if (v.position < position)  
            v.move(position - v.position);  
        else  
            move(10);  
    }  
};
```

```
class Truck : public Vehicle {  
public:  
    int load;  
    Truck() {}  
    void move(int d) {  
        if (d <= 55) position += d; else position += 55;  
    }  
};
```

```
#include <iostream>
using namespace std;

int main() {
    Truck t; // objects are stack-allocated
    Car c;
    c.passengers = 2;
    c.move(60);
    Vehicle *v = &c; // alias
    v->move(70);
    c.await(t);
}
```

On this example, object representation is not different from JAVA's

descr. Vehicle
position

descr. Car
position
passengers

descr. Truck
position
load

But in C++, we also **multiple inheritance**.

Consequence: we cannot use anymore the principle that

- the object layout for the super class is a prefix of the object layout of the subclass
- the descriptor for the super class is a prefix of the descriptor for the subclass

```
class Rocket {  
public:  
    float thrust;  
    Rocket() { }  
    virtual void display() {}  
};  
  
class RocketCar : public Car, public Rocket {  
public:  
    char *name;  
    void move(int d) { position += 2*d; }  
};
```

descr. RocketCar
position
passengers
descr. Rocket
thrust
name

Representations of Car and Rocket are appended.

In particular, a cast such as

```
RocketCar rc;  
... (Rocket)rc ...
```

is compiled using pointer arithmetic

```
... rc + 16 ...
```

This is not a no-op anymore.

descr. RocketCar
position
passengers
descr. Rocket
thrust
name

Let us now assume that Rocket also inherits from Vehicle

```
class Rocket : public Vehicle {  
public:  
    float thrust;  
    Rocket() { }  
    virtual void display() {}  
};  
  
class RocketCar : public Car, public Rocket {  
public:  
    char *name;  
    ...  
};
```

descr. RocketCar
position
passengers

descr. Rocket
position
thrust

name

We now have **two** fields position...

... and thus a possible ambiguity

```
class RocketCar : public Car, public Rocket {  
public:  
    char *name;  
    void move(int d) { position += 2*d; }  
};
```

```
vehicles.cc: In member function 'virtual void RocketCar::move(int)  
vehicles.cc:51:22: error: reference to 'position' is ambiguous
```

We have to say which one we refer to

```
class RocketCar : public Car, public Rocket {  
public:  
    char *name;  
    void move(int d) { Rocket::position += 2*d; }  
};
```

To have a single instance of `Vehicle` inside `RocketCar`, we need to modify the way `Car` and `Rocket` inherit from `Vehicle`;

This is **virtual inheritance**.


```
class Vehicle { ... };  
  
class Car : public virtual Vehicle { ... };  
  
class Rocket : public virtual Vehicle { ... };  
  
class RocketCar : public Car, public Rocket {
```

There is no ambiguity anymore:


```
public:  
    char *name;  
    void move(int d) { position += 2*d; }  
};
```

Three Class Diagrams


```
class Vehicle { ... };  
class Car : Vehicle { ... };  
class Rocket { ... };  
class RocketCar : Car, Rocket { ... };
```



```
class Vehicle { ... };  
class Car : Vehicle { ... };  
class Rocket : Vehicle { ... };  
class RocketCar : Car, Rocket { ... };
```



```
class Vehicle { ... };  
class Car : virtual Vehicle { ... };  
class Rocket : virtual Vehicle { ... };  
class RocketCar : Car, Rocket { ... };
```



g++'s command line option `-fdump-lang-class` outputs a text file containing objects and tables layout.

Though JAVA only features simple inheritance, interfaces make method call more complex, in a way analogous to multiple inheritance.

```
interface I {  
    void m();  
}  
  
class C {  
    void foo(I x) { x.m(); }  
}
```

When compiling `x.m()`, we have no idea what the class of object `x` will be.

Instead of dispatching according to the type of the object, we can use the types of **all** the actual parameters; this is called **multiple dispatch**.

An example: Julia, a mathematically-oriented language.

```
function +(x::Int64 , y::Int64 ) ... end
function +(x::Float64, y::Float64) ... end
function +(x::Date   , y::Time   ) ... end
```

Another example: CLOS (Common Lisp Object System).

Pattern matching, as we find in OCaml for instance, e.g.,

```
let rec eval = function
| Const n -> ...
| Call ("print", [e]) -> ...
| Call (f, el) -> ...
```

is a form of **dynamic dispatch**: the branch is selected according to some runtime information.